

US Accidents Exploratory Data Analysis

TODO - talk about EDA

TODO - talk about the dataset (source, what it contains, how it will be useful)

- Kaggle
- information about accidents
- can use useful to prevent accidents
- mention that this does not contain data about New York

```
In [2]: import opendatasets as od
```

```
download_url = 'https://www.kaggle.com/datasets/sobhanmoosavi/us-accidents'
```

```
od.download(download_url)
```

Skipping, found downloaded files in ".\us-accidents" (use force=True to force download)

```
In [3]: data_filename = './us-accidents/US_Accidents_March23.csv'
```

Data Preparation and Cleaning

1. Load the file using Pandas

2. Look at some information about the data & the columns

3. Fix any missing or incorrect values

```
In [4]: import pandas as pd
```

```
In [5]: df = pd.read_csv(data_filename)
```

```
In [6]: df
```

Out[6]:

	ID	Source	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	E
0	A-1	Source2	3	2016-02-08 05:46:00	2016-02-08 11:00:00	39.865147	-84.058723	NaN	
1	A-2	Source2	2	2016-02-08 06:07:59	2016-02-08 06:37:59	39.928059	-82.831184	NaN	
2	A-3	Source2	2	2016-02-08 06:49:27	2016-02-08 07:19:27	39.063148	-84.032608	NaN	
3	A-4	Source2	3	2016-02-08 07:23:34	2016-02-08 07:53:34	39.747753	-84.205582	NaN	
4	A-5	Source2	2	2016-02-08 07:39:07	2016-02-08 08:09:07	39.627781	-84.188354	NaN	
...	
7728389	A-7777757	Source1	2	2019-08-23 18:03:25	2019-08-23 18:32:01	34.002480	-117.379360	33.99888	-117.379360
7728390	A-7777758	Source1	2	2019-08-23 19:11:30	2019-08-23 19:38:23	32.766960	-117.148060	32.76555	-117.148060
7728391	A-7777759	Source1	2	2019-08-23 19:00:21	2019-08-23 19:28:49	33.775450	-117.847790	33.77740	-117.847790
7728392	A-7777760	Source1	2	2019-08-23 19:00:21	2019-08-23 19:29:42	33.992460	-118.403020	33.98311	-118.403020
7728393	A-7777761	Source1	2	2019-08-23 18:52:06	2019-08-23 19:21:31	34.133930	-117.230920	34.13736	-117.230920

7728394 rows × 46 columns

In [7]: df.columns

Out[7]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat', 'Start_Lng', 'End_Lat', 'End_Lng', 'Distance(mi)', 'Description', 'Street', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone', 'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Wind_Chill(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction', 'Wind_Speed(mph)', 'Precipitation(in)', 'Weather_Condition', 'Amenity', 'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight'], dtype='object')

In [8]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7728394 entries, 0 to 7728393
Data columns (total 46 columns):
#   Column                                Dtype
---  -
0   ID                                    object
1   Source                               object
2   Severity                             int64
3   Start_Time                           object
4   End_Time                             object
5   Start_Lat                            float64
6   Start_Lng                            float64
7   End_Lat                              float64
8   End_Lng                              float64
9   Distance(mi)                         float64
10  Description                           object
11  Street                                object
12  City                                  object
13  County                               object
14  State                                object
15  Zipcode                              object
16  Country                              object
17  Timezone                             object
18  Airport_Code                         object
19  Weather_Timestamp                    object
20  Temperature(F)                       float64
21  Wind_Chill(F)                        float64
22  Humidity(%)                          float64
23  Pressure(in)                         float64
24  Visibility(mi)                       float64
25  Wind_Direction                       object
26  Wind_Speed(mph)                      float64
27  Precipitation(in)                    float64
28  Weather_Condition                    object
29  Amenity                              bool
30  Bump                                  bool
31  Crossing                              bool
32  Give_Way                             bool
33  Junction                              bool
34  No_Exit                               bool
35  Railway                               bool
36  Roundabout                           bool
37  Station                               bool
38  Stop                                  bool
39  Traffic_Calming                      bool
40  Traffic_Signal                       bool
41  Turning_Loop                         bool
42  Sunrise_Sunset                       object
43  Civil_Twilight                       object
44  Nautical_Twilight                    object
45  Astronomical_Twilight                object
dtypes: bool(13), float64(12), int64(1), object(20)
memory usage: 2.0+ GB

```

```
In [9]: df.describe()
```

Out[9]:

	Severity	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Te
count	7.728394e+06	7.728394e+06	7.728394e+06	4.325632e+06	4.325632e+06	7.728394e+06	
mean	2.212384e+00	3.620119e+01	-9.470255e+01	3.626183e+01	-9.572557e+01	5.618423e-01	
std	4.875313e-01	5.076079e+00	1.739176e+01	5.272905e+00	1.810793e+01	1.776811e+00	
min	1.000000e+00	2.455480e+01	-1.246238e+02	2.456601e+01	-1.245457e+02	0.000000e+00	
25%	2.000000e+00	3.339963e+01	-1.172194e+02	3.346207e+01	-1.177543e+02	0.000000e+00	
50%	2.000000e+00	3.582397e+01	-8.776662e+01	3.618349e+01	-8.802789e+01	3.000000e-02	
75%	2.000000e+00	4.008496e+01	-8.035368e+01	4.017892e+01	-8.024709e+01	4.640000e-01	
max	4.000000e+00	4.900220e+01	-6.711317e+01	4.907500e+01	-6.710924e+01	4.417500e+02	

In [10]: df.describe(include = 'object')

Out[10]:

	ID	Source	Start_Time	End_Time	Description	Street	City	County	Stat
count	7728394	7728394	7728394	7728394	7728389	7717525	7728141	7728394	772839
unique	7728394	3	6131796	6705355	3761578	336306	13678	1871	4
top	A-1	Source1	2021-01-26 16:16:13	2021-11-22 08:00:00	A crash has occurred causing no to minimum del...	I-95 N	Miami	Los Angeles	C
freq	1	4325632	225	112	9593	78430	186917	526851	174143

In [11]: df.shape

Out[11]: (7728394, 46)

In [12]: numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']

numeric_df = df.select_dtypes(include=numerics)
len(numeric_df.columns)

Out[12]: 13

In [13]: missing_percentage = df.isna().sum().sort_values(ascending=False)/len(df)
missing_percentage

```
Out[13]: End_Lat          4.402935e-01
End_Lng          4.402935e-01
Precipitation(in) 2.851286e-01
Wind_Chill(F)     2.586590e-01
Wind_Speed(mph)   7.391355e-02
Visibility(mi)     2.291524e-02
Wind_Direction    2.267043e-02
Humidity(%)       2.253301e-02
Weather_Condition 2.244438e-02
Temperature(F)    2.120143e-02
Pressure(in)      1.820288e-02
Weather_Timestamp 1.555666e-02
Nautical_Twilight 3.007869e-03
Civil_Twilight    3.007869e-03
Sunrise_Sunset    3.007869e-03
Astronomical_Twilight 3.007869e-03
Airport_Code      2.928810e-03
Street            1.406372e-03
Timezone          1.010300e-03
Zipcode           2.477876e-04
City              3.273643e-05
Description        6.469649e-07
Traffic_Signal    0.000000e+00
Roundabout        0.000000e+00
Station           0.000000e+00
Stop              0.000000e+00
Traffic_Calming   0.000000e+00
Country           0.000000e+00
Turning_Loop      0.000000e+00
No_Exit           0.000000e+00
End_Time          0.000000e+00
Start_Time        0.000000e+00
Severity          0.000000e+00
Railway           0.000000e+00
Crossing          0.000000e+00
Junction          0.000000e+00
Give_Way          0.000000e+00
Bump              0.000000e+00
Amenity           0.000000e+00
Start_Lat         0.000000e+00
Start_Lng         0.000000e+00
Distance(mi)      0.000000e+00
Source            0.000000e+00
County            0.000000e+00
State             0.000000e+00
ID                0.000000e+00
dtype: float64
```

```
In [14]: missing_percentage[missing_percentage != 0]
```


```
Out[14]: End_Lat          4.402935e-01
End_Lng          4.402935e-01
Precipitation(in) 2.851286e-01
Wind_Chill(F)    2.586590e-01
Wind_Speed(mph)  7.391355e-02
Visibility(mi)    2.291524e-02
Wind_Direction    2.267043e-02
Humidity(%)       2.253301e-02
Weather_Condition 2.244438e-02
Temperature(F)    2.120143e-02
Pressure(in)      1.820288e-02
Weather_Timestamp 1.555666e-02
Nautical_Twilight 3.007869e-03
Civil_Twilight    3.007869e-03
Sunrise_Sunset    3.007869e-03
Astronomical_Twilight 3.007869e-03
Airport_Code      2.928810e-03
Street            1.406372e-03
Timezone          1.010300e-03
Zipcode           2.477876e-04
City              3.273643e-05
Description        6.469649e-07
dtype: float64
```

```
In [15]: type(missing_percentage)
```

```
Out[15]: pandas.core.series.Series
```

```
In [16]: missing_percentage[missing_percentage != 0].plot(kind='barh')
```

```
Out[16]: <Axes: >
```



- Remove columns that you don't want to use

Exploratory Data Analysis and Visualization

Columns we'll analyze

- City
- State
- Start Time
- Start Lat
- Start Lag
- Temperature
- Weather Condition

In [17]: `df.columns`

Out[17]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat', 'Start_Lng', 'End_Lat', 'End_Lng', 'Distance(mi)', 'Description', 'Street', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone', 'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Wind_Chill(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction', 'Wind_Speed(mph)', 'Precipitation(in)', 'Weather_Condition', 'Amenity', 'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight'], dtype='object')

Analysis for City

In [18]: `df.City`

Out[18]:

0	Dayton
1	Reynoldsburg
2	Williamsburg
3	Dayton
4	Dayton
...	
7728389	Riverside
7728390	San Diego
7728391	Orange
7728392	Culver City
7728393	Highland

Name: City, Length: 7728394, dtype: object

In [19]: `df.City.unique()`

Out[19]: array(['Dayton', 'Reynoldsburg', 'Williamsburg', ..., 'Ness City', 'Clarksdale', 'American Fork-Pleasant Grove'], dtype=object)

In [20]: `Unique_cities = df.City.unique()
len(Unique_cities)`

Out[20]: 13679

In [21]: `cities_by_accident = df.City.value_counts()
cities_by_accident`

```
Out[21]: Miami          186917
Houston        169609
Los Angeles    156491
Charlotte      138652
Dallas         130939
...
Benkelman      1
Old Appleton   1
Wildrose       1
Mc Nabb        1
American Fork-Pleasant Grove  1
Name: City, Length: 13678, dtype: int64
```

```
In [22]: 'New York' in df.City
```

```
Out[22]: False
```

```
In [23]: cities_by_accident[:20]
```


```
Out[23]: Miami          186917
Houston        169609
Los Angeles    156491
Charlotte      138652
Dallas         130939
Orlando        109733
Austin         97359
Raleigh        86079
Nashville      72930
Baton Rouge    71588
Atlanta        68186
Sacramento     66264
San Diego      55504
Phoenix        53974
Minneapolis    51488
Richmond       48845
Oklahoma City  46092
Jacksonville   42447
Tucson         39304
Columbia       38178
Name: City, dtype: int64
```

```
In [24]: type(cities_by_accident)
```

```
Out[24]: pandas.core.series.Series
```

```
In [25]: cities_by_accident[:20].plot(kind = 'barh')
```

```
Out[25]: <Axes: >
```

```
In [26]: import seaborn as sns
sns.set_style("darkgrid")
```

```
In [27]: sns.distplot(cities_by_accident)
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\3405282844.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(cities_by_accident)
<Axes: xlabel='City', ylabel='Density'>
```

Out[27]:



```
In [28]: # log_scale = True plot is more clear visible no of accident in cities  
sns.histplot(cities_by_accident, log_scale = True)
```

```
Out[28]: <Axes: xlabel='City', ylabel='Count'>
```



```
In [29]: cities_by_accident[cities_by_accident == 1]
```

```
Out[29]: Lake Andes          1
         Catoctin         1
         Duck Hill        1
         Westbrookville   1
         Saint Croix       1
         ..
         Benkelman        1
         Old Appleton     1
         Wildrose         1
         Mc Nabb          1
         American Fork-Pleasant Grove 1
         Name: City, Length: 1023, dtype: int64
```

```
In [30]: high_accident_cities = cities_by_accident[cities_by_accident>=1000]
         low_accident_cities = cities_by_accident[cities_by_accident<1000]
```

```
In [31]: len(high_accident_cities)/ len(cities_by_accident)
```

```
Out[31]: 0.08904810644831115
```

- 8.9% cities high accidents

```
In [32]: len(low_accident_cities)/len(cities_by_accident)
```

```
Out[32]: 0.9109518935516888
```

- 91% cities low accidents

```
In [33]: sns.distplot(high_accident_cities)
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\2843252471.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(high_accident_cities)
```

```
Out[33]: <Axes: xlabel='City', ylabel='Density'>
```



```
In [34]: sns.distplot(low_accident_cities)
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\469555131.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(low_accident_cities)
```

```
Out[34]: <Axes: xlabel='City', ylabel='Density'>
```



Analysis for Start Time

```
In [35]: df.Start_Time
```

```
Out[35]: 0      2016-02-08 05:46:00
1      2016-02-08 06:07:59
2      2016-02-08 06:49:27
3      2016-02-08 07:23:34
4      2016-02-08 07:39:07
...
7728389 2019-08-23 18:03:25
7728390 2019-08-23 19:11:30
7728391 2019-08-23 19:00:21
7728392 2019-08-23 19:00:21
7728393 2019-08-23 18:52:06
Name: Start_Time, Length: 7728394, dtype: object
```

```
In [36]: df.Start_Time[0]
```

```
Out[36]: '2016-02-08 05:46:00'
```

```
In [37]: df['Start_Time'] = pd.to_datetime(df['Start_Time'])
```

```
In [38]: df.dtypes
```

```
Out[38]: ID                object
Source                object
Severity              int64
Start_Time            datetime64[ns]
End_Time              object
Start_Lat             float64
Start_Lng             float64
End_Lat               float64
End_Lng               float64
Distance(mi)          float64
Description            object
Street                object
City                  object
County                object
State                 object
Zipcode               object
Country               object
Timezone              object
Airport_Code          object
Weather_Timestamp     object
Temperature(F)        float64
Wind_Chill(F)         float64
Humidity(%)           float64
Pressure(in)          float64
Visibility(mi)        float64
Wind_Direction        object
Wind_Speed(mph)       float64
Precipitation(in)     float64
Weather_Condition     object
Amenity               bool
Bump                  bool
Crossing              bool
Give_Way              bool
Junction              bool
No_Exit               bool
Railway               bool
Roundabout            bool
Station               bool
Stop                  bool
Traffic_Calming       bool
Traffic_Signal        bool
Turning_Loop          bool
Sunrise_Sunset        object
Civil_Twilight         object
Nautical_Twilight     object
Astronomical_Twilight object
dtype: object
```

```
In [39]: df.Start_Time[0]
```

```
Out[39]: Timestamp('2016-02-08 05:46:00')
```

```
In [40]:  #(norm_hist = True use to convert into percentage)
sns.distplot(df.Start_Time.dt.hour, bins = 24, kde = False, norm_hist = True , color
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\823610032.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Start_Time.dt.hour, bins = 24, kde = False, norm_hist = True, color = 'red')
```

Out[40]: <Axes: xlabel='Start_Time'>



What time of the day are accidents most frequent in?

- A high percentage of accidents occur between 7 am to 8 am (probably people in a hurry to get to work)
- Next highest percentage is 3 pm to 5 pm.

In [41]: `sns.distplot(df.Start_Time.dt.dayofweek, bins = 7, kde = False, norm_hist = True, color = 'red')`

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\530816872.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Start_Time.dt.dayofweek, bins = 7, kde = False, norm_hist = True, color = 'blue')
```

Out[41]: <Axes: xlabel='Start_Time'>



Is distribution of accidents by hour the same on weekends as on weekdays

```
In [42]: # dayofweek == 6 means sunday
sundays_start_time = df.Start_Time[df.Start_Time.dt.dayofweek == 6]
sns.distplot(sundays_start_time.dt.hour, bins = 24, kde = False, norm_hist = True)
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\2255068480.py:3: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(sundays_start_time.dt.hour, bins = 24, kde = False, norm_hist = True)
```

Out[42]: <Axes: xlabel='Start_Time'>



```
In [43]: mondays_start_time = df.Start_Time[df.Start_Time.dt.dayofweek == 0]
sns.distplot(mondays_start_time.dt.hour, bins = 24, kde = False, norm_hist = True)
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\1425569934.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(mondays_start_time.dt.hour, bins = 24, kde = False, norm_hist = True)
```

```
Out[43]: <Axes: xlabel='Start_Time'>
```



- **On workdays i.e. Monday, Tuesday, Wednesday, Thursday, Friday you'll find almost the same trend in accident time.
- While on Saturday and Sunday there is a different trend i.e. from 10 am to 7 pm the frequency of accident is more.**

Analysis for Month distribution

```
In [44]: sns.distplot(df.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'green')
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\1242160542.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'green')
```

```
Out[44]: <Axes: xlabel='Start_Time'>
```



- The accidents are high from December and it is lowest at July. The rise continues to increase from the month of July.
- It's seems during summer there are less accidents but as the winter starts the is a increasing trend in accidents.

Analysis for Year

```
In [45]: df.Start_Time.dt.year
```

```
Out[45]: 0      2016
1      2016
2      2016
3      2016
4      2016
...
7728389 2019
7728390 2019
7728391 2019
7728392 2019
7728393 2019
Name: Start_Time, Length: 7728394, dtype: int64
```

```
In [46]: df_2019 = df[df.Start_Time.dt.year == 2019]
sns.distplot(df_2019.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True,
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\3765400231.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_2019.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'Blue')
```

<Axes: xlabel='Start_Time'>

Out[46]:



In [47]:

```
df_2020 = df[df.Start_Time.dt.year == 2020]
sns.distplot(df_2020.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True,
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\2660523566.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_2020.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'Blue')
```

<Axes: xlabel='Start_Time'>

Out[47]:



```
In [48]: df_2021 = df[df.Start_Time.dt.year == 2021]
sns.distplot(df_2021.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True,
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\607408822.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_2021.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'Blue')
```

```
Out[48]: <Axes: xlabel='Start_Time'>
```



```
In [49]: df_2022 = df[df.Start_Time.dt.year == 2022]
sns.distplot(df_2022.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True,
```

C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\1463311549.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_2022.Start_Time.dt.month, bins = 12, kde = False, norm_hist = True, color = 'Blue')
```

```
Out[49]: <Axes: xlabel='Start_Time'>
```




- Much data is missing for yearly analysis
- so, need to check some other column affected by our study, we can analysis source dataset

In [50]: `df.Source`

```
Out[50]:
0      Source2
1      Source2
2      Source2
3      Source2
4      Source2
...
7728389 Source1
7728390 Source1
7728391 Source1
7728392 Source1
7728393 Source1
Name: Source, Length: 7728394, dtype: object
```


```
In [51]: df_2019 = df[df.Start_Time.dt.year == 2019]
df_2019_Source1=df_2019[df_2019.Source == 'Source1']
#sns.distplot(df_2019_Source1.Start_Time.dt.month, bins=12, kde=False,norm_hist = 1
sns.histplot(df_2019_Source1['Start_Time'].dt.month, color='blue', bins=12, stat='p
```

```
Out[51]: <Axes: xlabel='Start_Time', ylabel='Percent'>
```



```
In [52]: df_2019 = df[df.Start_Time.dt.year == 2019]
df_2019_Source2=df_2019[df_2019.Source == 'Source2']
#sns.distplot(df_2019_Source2.Start_Time.dt.month, bins=12, kde=False,norm_hist=True)
sns.histplot(df_2019_Source2['Start_Time'].dt.month, color='blue', bins=12, stat='p
```

```
Out[52]: <Axes: xlabel='Start_Time', ylabel='Percent'>
```




```
In [53]: df_2019 = df[df.Start_Time.dt.year == 2019]
df_2019_Source3=df_2019[df_2019.Source == 'Source3']
```



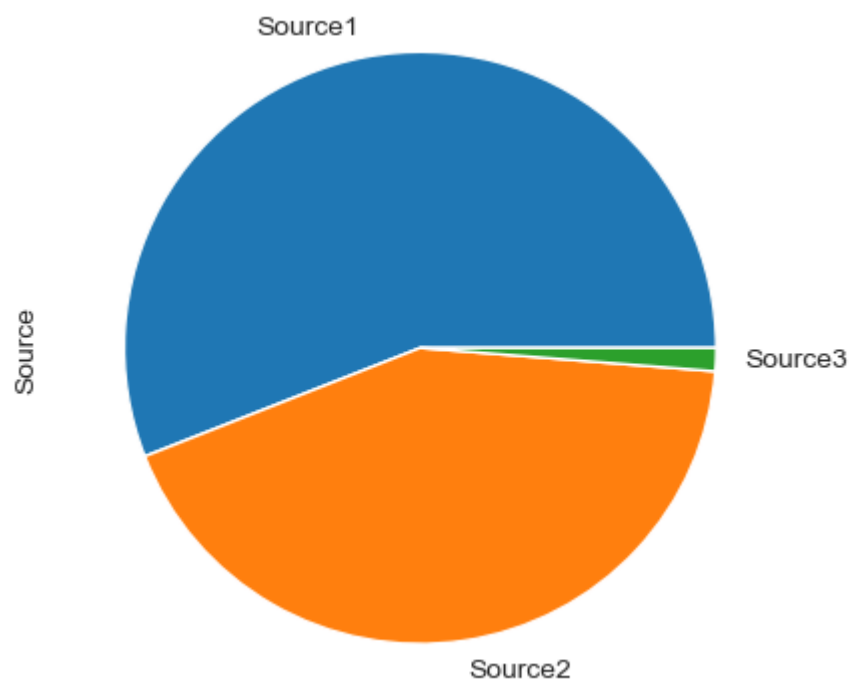
```
#sns.distplot(df_2019_Source3.Start_Time.dt.month, bins=12, kde=False, stat='percent')  
sns.histplot(df_2019_Source3['Start_Time'].dt.month, color='blue', bins=12, stat='percent')
```

Out[53]: <Axes: xlabel='Start_Time', ylabel='Percent'>



In [54]: df.Source.value_counts().plot(kind = 'pie')

Out[54]: <Axes: ylabel='Source'>



- There seems to be some issue with the Source2 and Source3 data so consider excluding Source2 and Source3 **

Start Latitude & Longitude

```
In [55]: df.Start_Lat
```

```
Out[55]: 0      39.865147
1      39.928059
2      39.063148
3      39.747753
4      39.627781
...
7728389 34.002480
7728390 32.766960
7728391 33.775450
7728392 33.992460
7728393 34.133930
Name: Start_Lat, Length: 7728394, dtype: float64
```

```
In [56]: df.Start_Lng
```

```
Out[56]: 0      -84.058723
1      -82.831184
2      -84.032608
3      -84.205582
4      -84.188354
...
7728389 -117.379360
7728390 -117.148060
7728391 -117.847790
7728392 -118.403020
7728393 -117.230920
Name: Start_Lng, Length: 7728394, dtype: float64
```

```
In [57]: # use sample function to extract 10% Data
sample_df = df.sample(int(0.1 * len(df)))
```

```
In [58]: sns.scatterplot(y = sample_df.Start_Lat, x = sample_df.Start_Lng, size = 0.001)
```

```
Out[58]: <Axes: xlabel='Start_Lng', ylabel='Start_Lat'>
```



```
In [59]: # show the above Lat & Lng scatter plot in Map (use Libraries folium)
import folium
```

```
In [60]: lat,lon = df.Start_Lng[0],df.Start_Lat[0]
lat,lon
```

```
Out[60]: (-84.058723, 39.865147)
```

```
In [61]: # sample().iteritems() used to show only 100 results
for x in df[['Start_Lat','Start_Lng']].sample(100).iteritems():
    print(x[1])
```

```
7010944    40.012440
1321647    30.212299
821239     36.052238
4066706    39.872330
104620     33.862186
```

```
...
```

```
1477305    39.890221
3986384    34.947461
3432428    39.105720
6666424    36.200835
3630864    45.101930
```

```
Name: Start_Lat, Length: 100, dtype: float64
```

```
7010944    -77.537700
1321647    -82.639572
821239     -86.743103
4066706    -75.348974
104620     -118.041985
```

```
...
```

```
1477305    -76.237991
3986384    -89.833700
3432428    -94.840850
6666424    -86.769689
3630864    -93.456630
```

```
Name: Start_Lng, Length: 100, dtype: float64
```

```
C:\Users\Atul Gupta\AppData\Local\Temp\ipykernel_11296\3450368729.py:2: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.
  for x in df[['Start_Lat', 'Start_Lng']].sample(100).iteritems():
```

```
In [62]: # creat heatmap
from folium.plugins import HeatMap
```


```
In [63]: # zip used to pair the both list, it is nescessary to convert list of lat & lng pair
zip(list(df.Start_Lat), list(df.Start_Lng))
```

```
Out[63]: <zip at 0x1c3e1ea7e00>
```

```
In [64]: sample_df = df.sample(int(0.001*len(df)))
lat_lng_pairs = zip(list(df.Start_Lat), list(df.Start_Lng))
```

```
In [65]: map = folium.Map()
HeatMap(lat_lng_pairs).add_to(map)
map
```

```
Out[65]: Make this Notebook Trusted to load map: File -> Trust Notebook
```



Are there more accidents in warmer or colder areas

```
In [66]: df['Temperature(F)']
```

```
Out[66]:
0      36.9
1      37.9
2      36.0
3      35.1
4      36.0
...
7728389  86.0
7728390  70.0
7728391  73.0
7728392  71.0
7728393  79.0
Name: Temperature(F), Length: 7728394, dtype: float64
```

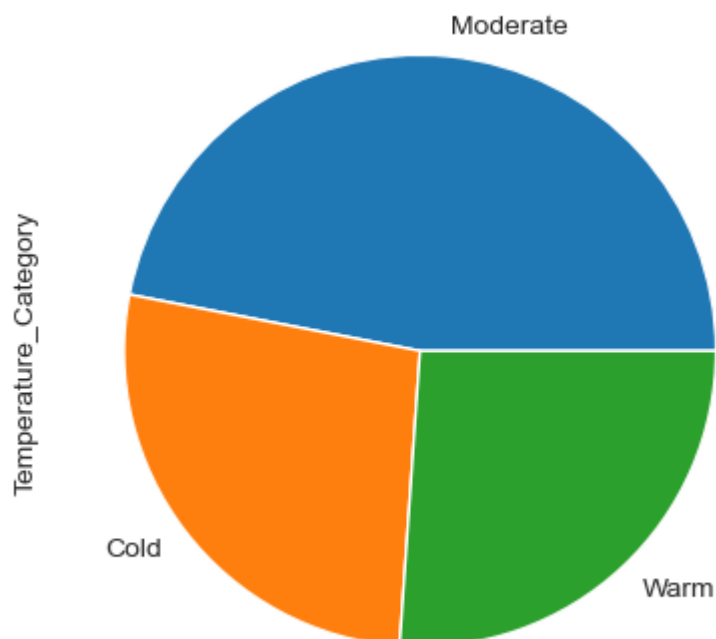
```
In [67]: # Create temperature bins (customize according to your data)
bins = [0, 50, 75, 100]
labels = ['Cold', 'Moderate', 'Warm']
# Assign temperature ranges to each row
df['Temperature_Category'] = pd.cut(df['Temperature(F)'], bins=bins, labels=labels,
df['Temperature_Category'])
```

```
Out[67]: 0          Cold
1          Cold
2          Cold
3          Cold
4          Cold
...
7728389    Warm
7728390    Moderate
7728391    Moderate
7728392    Moderate
7728393    Warm
Name: Temperature_Category, Length: 7728394, dtype: category
Categories (3, object): ['Cold' < 'Moderate' < 'Warm']
```

```
In [68]: # Group by temperature category and calculate the number of accidents
#accidents_by_temperature = df.groupby('Temperature_Category', observed=False).size

accidents_by_temperature = df['Temperature_Category'].value_counts()
df['Temperature_Category'].value_counts().plot(kind='pie')
```

```
Out[68]: <Axes: ylabel='Temperature_Category'>
```



- Moderate temperature days more accident happens


Analyzing the data by state

```
In [69]: states = df['State'].value_counts().head(5)
states
# The data indicates california is the highest accident state
```

```
Out[69]: CA    1741433  
FL     880192  
TX     582837  
SC     382557  
NY     347960  
Name: State, dtype: int64
```

```
In [70]: sns.barplot(y=states , x = states.index, palette="RdPu")
```

```
Out[70]: <Axes: ylabel='State'>
```



Summary and Conclusion

Insights:

- The cities with the highest reported accidents are Miami, Houston, Los Angeles, Charlotte, Dallas, Orlando, Austin, Raleigh, Nashville, Baton Rouge, Atlanta, Sacramento, San Diego, Phoenix, Minneapolis, Richmond, Oklahoma City, Jacksonville, Tucson, and Columbia.
- About 8.9% of cities experience a high number of accidents.
- The majority of cities (91%) have a low number of accidents.
- Over 1023 cities reported just 1 accident, suggesting the presence of potential outliers that may need to be addressed.
- There is a notable spike in accidents around 7-8 AM, possibly correlated with morning rush hours and commuting to work or school.
- Another spike occurs around 4-5 PM, likely associated with evening rush hours and the return home from work or recreational activities.
- On weekdays (Monday to Friday), the trend in accident times is consistent.
- On weekends (Saturday and Sunday), there is a different trend, with a higher frequency of accidents between 10 AM and 7 PM.

- There is a seasonal variation in accidents, with fewer incidents during the summer and an increasing trend as winter approaches.
- The use of Folium indicates that many people live near bay areas.
- No data from New York
- California, Florida, Texas, South Carolina, and New York emerge as the top 5 states with the highest number of accidents
- Moderate temperature days more accident happens

In []: